Arcolens

Snapshot: September 27, 2026

This report contains candidate insights, chart configurations, confidence and bias notes, and an audit of existing published statements. It also records the data-quality and website issues uncovered during the work, and the status of the recurring process.

Review at a glance

Input. 1,698 incident records in the supplied Arcola CSV. Figures describe this snapshot; the public tracker currently uses different classifications.

Candidate insights. 13 candidates: five audited starting examples, six findings selected through discovery, and two descriptive profiles for potential causes and national security. Each appears below with wording, settings, and limitations.

Published insights. 19 current empirical statement blocks: six supported, two needing updates, one unsupported under its stated interpretation, and ten unresolved. Unresolved includes ambiguous definitions and missing evidence; it is not automatically a finding of falsehood.

Additional findings. A misleading severity headline, an ineffective severity preset, a source-scope problem in record 1607, a mismatch between the live mirror and export, and specific definition and validation gaps.

Recurring process. Runnable calculations, export comparison, model drafting, independent review, and a local monthly check. An unchanged-input rehearsal passed. A genuine legacy-export comparison is now complete; a second export verified from the current Arcola pipeline remains needed for a same-pipeline update.

How to interpret the findings

These are descriptive results from a selectively reported collection. Record counts are not necessarily counts of distinct events, reports, or people. No exposure denominator is available. Changes can reflect reporting, coverage, dating, and classification as well as underlying incidents. A falling share does not establish improved AI safety.

Calculation confidence and interpretation confidence are separate. Reproduced arithmetic does not validate model-assigned labels. Annual comparisons exclude incomplete 2026. Direct, indirect, and inferred severity require explicit channel and dimension definitions. Inferred harm does not establish that harm occurred.

The dashboard provides item-level calculation tables and full source records. This report includes the actual candidate wording, chart instructions, limitations, and claim assessments so it can be reviewed independently of the dashboard. Nothing has been published or sent to Simon.

Timeline and annual totals

IT-C002 - Capability / robustness: share down, count up

Among incident records dated 2021-2025 in the September 27, 2026 export, the share whose primary subdomain is "7.3 Lack of capability or robustness" fell each year, from 33 of 80 (41.25%) to 45 of 452 (9.96%). The count rose from 33 to 45. The falling share describes database composition and does not establish improved AI robustness.

Chart. Timeline by classification. Exact-year preset needed.

Settings. January 1, 2021-December 31, 2025; exclude incomplete 2026.; Proportions of each year's full incident-record cohort; switch to Counts to show the numerator increasing from 33 to 45.; stack or group by Primary Subdomain; stackBy=subdomain; all legend series visible.. Filters: No category filters and no hidden series. Host preset: year=['2021','2022','2023','2024','2025']; viewMode=proportion. Do not filter to 7.3.

Compatibility. An exact-year host preset is required. All Subdomains proportions use assignment denominators and do not reproduce any-label incident percentages. Quoted values require matching snapshot data. The exact-year host preset is supported by source inspection but has not been verified end to end here; use the annual supporting table until tested with matching data.

Confidence. Calculation: High for primary/any-label counting and exact annual comparisons in the export. Interpretation: Moderate for a compositional database statement; low for any claim that robustness has improved or failures have become rarer.

- 2021 is a local high point in the longer series. The finding is specific to 2021-2025.

- The category count rose from 33 to 45; its share fell because the rest of the database grew faster.

- The partial 2026 cohort differs from 2025. It is excluded from these complete-year comparisons.

Year	Selected records	All dated records	Share
2021	33	80	41.25%
2022	33	108	30.56%
2023	37	179	20.67%
2024	44	304	14.47%
2025	45	452	9.96%

IT-D006 - False or misleading information grew from a small base

In the September 27, 2026 export, records whose primary subdomain is "3.1 False or misleading information" accounted for 1 of 80 (1.25%) dated 2021, 9 of 108 (8.33%) in 2022, 21 of 179 (11.73%) in 2023, 32 of 304 (10.53%) in 2024, and 63 of 452 (13.94%) in 2025. The share did not rise every year, and the starting numerator is one record.

Chart. Timeline by classification. Exact-year preset needed.

Settings. January 1, 2021-December 31, 2025; exclude incomplete 2026.; Proportions of each year's full incident-record cohort; switch to Counts to distinguish count increases from share changes.; stack or group by Primary Subdomain; stackBy=subdomain; all legend series visible.. Filters: No category filters or hidden series. Host preset: year=['2021','2022','2023','2024','2025']; viewMode=proportion. Do not filter to 3.1.

Compatibility. An exact-year host preset is required because ordinary Year >= controls cannot exclude 2026. All Subdomains proportions use assignment denominators. Quoted values require matching snapshot data. The exact-year host preset is supported by source inspection but has not been verified end to end here; use the annual supporting table until tested with matching data.

Confidence. Calculation: High for the exported classifications and exact annual comparisons in this export. Interpretation: Moderate as a description of database composition; low for any claim about real-world incidence, risk or safety.

- The 2021 numerator is just one record: this is an exploratory finding from a small base.

- The share did not rise every year: it fell between 2023 and 2024.

- Reporting, classification and coverage changes remain alternative explanations.

Year	Selected records	All dated records	Share
2021	1	80	1.25%
2022	9	108	8.33%
2023	21	179	11.73%
2024	32	304	10.53%
2025	63	452	13.94%

IT-C001 - More records in each complete year

In the September 27, 2026 export, the number of incident records dated in each year rose from 80 (2021) to 108 (2022), 179 (2023), 304 (2024), and 452 (2025). These are record counts; growth can reflect reporting and coverage as well as underlying incidents.

Chart. Annual incident totals, Timeline by classification. Exact-year preset needed.

Settings. January 1, 2021-December 31, 2025; September 27, 2026 export.; Counts covering all records in each annual cohort, or 100% of that cohort; sum the domain segments.; stack or group by Domain; stackBy=domain; all seven domains visible.. Filters: No category filters. Host preset: year=['2021','2022','2023','2024','2025']; viewMode=counts.

Compatibility. Ordinary Year >= controls cannot close the interval at 2025. An exact-year host preset and matching snapshot data are required; otherwise use a supporting table. The exact-year host preset is supported by source inspection but has not been verified end to end here; use the annual supporting table until tested with matching data.

Confidence. Calculation: High for the retained distinct-record counts and date parsing; not an estimate of new submissions, events or people. Interpretation: Low for actual incident incidence or safety changes. Moderate as a description of this database snapshot.

- Records are not distinct events, reports or people.

- The choice of 2021 matters: 2020 has 92 records, above 2021. Changes in reporting and coverage can change the series.

- Historical backfill cannot be measured reliably from this single export. The ordinary ID assignment rule is documented, but IDs do not establish historical entry dates or past classifications.

Year	Records
2021	80
2022	108
2023	179
2024	304
2025	452

Overview pie

IT-C003 - Deepfake video leads the single-purpose tags

In the September 27, 2026 export, the most frequent single primary AI-purpose tag is "Deepfake Video Generation" (316 of 1,698 records, 18.6%), followed by "Chatbot" (181, 10.7%). This ranking covers all dates and describes primary tags; including secondary purposes changes the order below first place.

Chart. Overview pie. Chart view available.

Settings. Observed record dates September 26, 1983-September 25, 2026; September 27, 2026 snapshot. Partial 2026 is included only in snapshot composition.; Pie proportions with counts in tooltips; denominator is all 1,698 export records.; stack or group by AI Purpose, primary; stackBy=aiGoal.. Filters: No filters; retain all dates and primary-purpose scope.

Compatibility. C03 supports this primary-purpose composition. Other aggregates remaining purposes. Use a supporting table for frozen-export values until chart data match the export.

Confidence. Calculation: High for the primary-purpose ranking; alternate purposes are counted separately. Interpretation: Moderate for the retained tag ranking; low for which AI applications are intrinsically riskiest.

- The ranking depends on category breadth: related assistant uses are split across labels.
- The grouped Other slice combines the remaining labels and is not a single competing purpose.
- These classifications do not establish which applications are intrinsically most risky.

Within risk domains

IT-C004 - Fraud within the malicious-actor domain

Of the incident records dated 2016-2025 whose primary domain is "4 Malicious actors," 368 of 535 (68.8%) have primary subdomain "4.3 Fraud, scams and targeted manipulation" in the September 27, 2026 export. This is a share within the selected domain, not within the entire database.

Chart. Overview domain / subdomain bars, Overview pie. Chart view available.

Settings. January 1, 2016-December 31, 2025. Supporting pooled periods: 2016-2023 and 2024-2025.; Counts by primary subdomain; accompanying share is 368 of 535 domain records (68.8%).; stack or group by Subdomain (the primary subdomain field); stackBy=subdomain; all subdomain segments visible.. Filters: Domain='4 Malicious actors'; Year selections: 2016, 2017, 2018, 2019, 2020, 2021, 2022, 2023, 2024, 2025.

Compatibility. C04 supports exact Year selections and Subdomain counts (primary field). It has no visible Proportions toggle; show supplied percentages in accompanying text or a table. Chart data must match the export.

Confidence. Calculation: High for the explicitly primary, closed-window numerator and denominator. Interpretation: Moderate for category composition; low for the prevalence or absolute frequency of real-world malicious AI use.

- The denominator is 535 primary malicious-actor records, not all 1,698 records.
- The pooled window is dominated by recent years: 413 of the 535 records are dated 2024-2025.
- Voluntary incident reports cannot establish the prevalence of malicious AI use.

IT-D001 - False or misleading information within misinformation

Among all records dated through December 31, 2025 whose primary domain is "3 Misinformation," 128 of 141 (90.8%) have primary subdomain "3.1 False or misleading information" in the September 27, 2026 export. The bar chart's 2016-2025 window gives 127 of 140 (90.7%).

Chart. Overview domain / subdomain bars. Chart window differs.

Settings. Headline: all records through December 31, 2025, observed April 5, 2011-December 31, 2025. C04: January 1, 2016-December 31, 2025.; Headline: 128/141 (90.8%). C04 window: 127/140 (90.7%). Display bar counts and put these supplied shares in accompanying text or a table.; stack or group by Subdomain (the primary subdomain field);

stackBy=subdomain; all subdomain segments visible.. Filters: C04 Domain='3 Misinformation'; Year selections: 2016, 2017, 2018, 2019, 2020, 2021, 2022, 2023, 2024, 2025.

Compatibility. C04 cannot show pre-2016 records and has no visible Proportions toggle. Its narrower result cannot substitute for the full-window headline. Frozen-export values require matching data.

Confidence. Calculation: High for the explicitly primary numerator and denominator. Interpretation: Moderate for category composition; low for the prevalence of real-world events.

- The denominator is the primary misinformation domain, not the full dataset.

- The headline includes complete years before 2016; the bar chart starts at 2016. Match the window before publishing.

- The small taxonomy within this domain partly determines the concentration.

IT-D005 - Overreliance within human-computer interaction

Of the 60 records dated through 2025 whose primary domain is "5 Human-Computer Interaction," 53 (88.3%) have primary subdomain "5.1 Overreliance and unsafe use" in the September 27, 2026 export. Restricted to 2016-2025, the bar chart's window, the figures are 52 of 56 (92.9%). Four earlier records account for the difference between those windows.

Chart. Overview domain / subdomain bars. Chart window differs.

Settings. Headline: all records through December 31, 2025, observed December 10, 2012-November 10, 2025. C04: January 1, 2016-December 31, 2025.; Headline: 53/60 (88.3%). C04 window: 52/56 (92.9%). Display bar counts; place supplied shares in accompanying text or a table.; stack or group by Subdomain (the primary subdomain field); stackBy=subdomain; all subdomain segments visible.. Filters: C04 Domain='5 Human-Computer Interaction'; Year selections: 2016, 2017, 2018, 2019, 2020, 2021, 2022, 2023, 2024, 2025.

Compatibility. C04 cannot show pre-2016 records and has no visible Proportions toggle. Use the supporting table for the full-window headline and matching snapshot data for quoted chart values.

Confidence. Calculation: High for the explicitly primary numerator and denominator. Interpretation: Moderate for category composition; low for the prevalence of real-world events.

- The denominator is only 60 primary human-computer-interaction records.

- The headline covers all records through 2025; the bar chart starts at 2016. Four earlier records explain the window difference.

- A suspected duplicate pair warrants source review. Excluding record 353 as a sensitivity gives 52 of 59 (88.1%); the supplied export is unchanged.

Severity and evidence

IT-C005 - Severity counts change when inference is included

In the September 27, 2026 export, 148 of 1,698 records (8.7%) have a direct or indirect severity score of at least 4 in an original harm dimension; including model-inferred scores gives 286 (16.8%). Both counts use the original ten dimensions and do not count zero-coded cells toward the threshold. Inferred scores do not establish that the harms occurred.

Chart. Incident severity detail, Overview pie, Supporting table. Supporting table needed.

Settings. Observed record dates September 26, 1983-September 25, 2026; September 27, 2026 snapshot, including partial 2026.; Supporting rows over 1,698 records: ten dimensions, direct only 90 (5.3%); direct or indirect 148 (8.7%); all three channels 286 (16.8%). Thirteen dimensions, all channels: 376 (22.1%).; stack or group by Unavailable: no current aggregate grouping separates these channel maxima and dimension sets.. Filters: Table definition: all records; score >=4 in at least one dimension; lower_bound missing-score policy. State channels and dimension count separately for each row.

Compatibility. No aggregate control reproduces the ten-dimension direct/indirect comparison. The public headline includes all three channels across thirteen dimensions. C05 inspects one incident and omits an inferred numeric series.

Confidence. Calculation: High for the observed score thresholds and channel maxima. The 2 zero-coded records (#413, #1607) and 0 blank-cell records have explicit bounds and an exclusion sensitivity; whether a zero score means 'not assessed' or 'no harm' is unresolved with the exporter. Interpretation: High that channel choice changes the retained score count; low that inferred scores verify realized harm severity.

- Inferred scores are not confirmation that the harms occurred. Zero-coded cells do not count toward the threshold; the exporter has not confirmed their meaning.

- This comparison uses the original ten harm dimensions. The three additions - psychological harm, epistemic harm, and CSEA - lack category-specific threshold rubrics in the public sources examined.

- Use a supporting table. The public headline includes inferred scores across thirteen dimensions, and its current direct-only label is misleading.

Publication note. Show this comparison in a labeled table. The public headline combines thirteen dimensions and includes inferred scores, so it represents a different measure.

IT-D002 - The share meeting a severity threshold fell

In the September 27, 2026 export, the share of incident records with a direct or indirect severity score of at least 3 in an original harm dimension fell each year from 35 of 80 (43.8%) in 2021 to 90 of 452 (19.9%) in 2025, while the count rose from 35 to 90. The measure uses the original ten dimensions, excludes inferred scores, and does not count zero-coded cells toward the threshold. It describes database composition, not changes in underlying safety.

Chart. Supporting table. Supporting table needed.

Settings. January 1, 2021-December 31, 2025; exclude incomplete 2026.; Show qualifying count, full annual incident-record denominator, and percentage for each year.; stack or group by Unavailable: no aggregate grouping matches this definition.. Filters: Table definition: all records; direct or indirect score >=3 in at least one of the original ten dimensions; lower_bound missing-score policy; inferred scores excluded.

Compatibility. Supporting table only. The public headline includes inferred scores across thirteen dimensions and cannot reproduce this ten-dimension direct/indirect measure.

Confidence. Calculation: High for the exported classifications and exact annual comparisons in this export. Interpretation: Moderate as a description of database composition; low for any claim about real-world incidence, risk or safety.

- The count meeting the threshold rose from 35 to 90 while its share fell.

- This depends on the original ten dimensions, direct/indirect channels and threshold 3.

- The present aggregate chart cannot faithfully reproduce this channel-specific calculation.

Year	Selected records	All dated records	Share
2021	35	80	43.75%
2022	36	108	33.33%
2023	53	179	29.61%
2024	86	304	28.29%
2025	90	452	19.91%

Additional table findings

IT-D003 - A larger share tagged as tool-type systems

In the September 27, 2026 export, records with the behavior tag "Tool" accounted for 19 of 80 (23.8%) dated 2021, 36 of 108 (33.3%) in 2022, 96 of 179 (53.6%) in 2023, 171 of 304 (56.3%) in 2024, and 256 of 452 (56.6%) in 2025. The share increased between the endpoints; the final two years were similar. Denominators include other and unknown tags.

Chart. Supporting table. Supporting table needed.

Settings. January 1, 2021-December 31, 2025; exclude incomplete 2026.; Show Tool count, all-record annual denominator, and percentage; retain all other behavior tags, including unknown in the denominator.; stack or group by Not applicable: behavior type has no current chart grouping.. Filters: No population exclusions; target behaviour_type=Tool.

Compatibility. Supporting table only; behavior type is not an available chart grouping. No legend or Proportions chart control applies to this table.

Confidence. Calculation: High for the exported classifications and exact annual comparisons in this export. Interpretation: Moderate as a description of database composition; low for any claim about real-world incidence, risk or safety.

- Tool is an LLM-assigned behavior tag whose meaning needs confirmation.

- The public chart has no matching grouping; use a labeled supporting table.

- The result describes database composition, not a population trend in AI deployment.

Publication note. Confirm the classification rule for "Tool" before selecting this finding for publication.

Year	Selected records	All dated records	Share
2021	19	80	23.75%
2022	36	108	33.33%
2023	96	179	53.63%
2024	171	304	56.25%
2025	256	452	56.64%

IT-D004 - A larger share tagged as human-controlled

In the September 27, 2026 export, the share of incident records tagged "Human-controlled" in the national-security autonomy field rose each year, from 19 of 80 (23.8%) dated 2021 to 349 of 452 (77.2%) dated 2025. The denominator includes every record, not only national-security incidents. This describes assigned tags and does not establish a trend in deployment.

Chart. Supporting table. Supporting table needed.

Settings. January 1, 2021-December 31, 2025; exclude incomplete 2026.; Show Human-controlled count, full annual incident-record denominator, and percentage.; stack or group by Not applicable: this autonomy field has no current chart grouping.. Filters: No population exclusions or national-security incident filter; target natsec_autonomy_level=Human-controlled.

Compatibility. Supporting table only; this autonomy field is not an available chart grouping.

Confidence. Calculation: High for the exported classifications and exact annual comparisons in this export. Interpretation: Moderate as a description of database composition; low for any claim about real-world incidence, risk or safety.

- The public framework defines Human-controlled as AI assisting while humans make final decisions. The tag's accuracy in individual records has not been established.

- The national-security autonomy field is populated across this export. Its application outside national-security incidents requires care.

- Use a supporting table; the current public chart has no grouping for this field.

Year	Selected records	All dated records	Share
2021	19	80	23.75%
2022	42	108	38.89%
2023	113	179	63.13%
2024	201	304	66.12%
2025	349	452	77.21%

Potential causes and national security

IT-P001 - Potential-cause annotations span six categories

In the 2026-09-27 export, 1,220 of 1,698 records (71.8%) have at least one non-placeholder potential-cause annotation in each of the six diagram categories. These are model-proposed explanations in a prompted framework, not six independently established causes per incident.

Chart. Potential-causes diagram. Diagram example and supporting table.

Settings. All dates in the supplied export; no time comparison; Potential Causes in the incident detail; counts and proportions do not apply to the diagram; stack or group by None; six qualitative branches. Filters: incident_id: 1

Compatibility. The diagram can illustrate record #1; the aggregate six-category coverage requires the attached table. Ignore N/A placeholders. The live mirror may contain different annotations.

Confidence. Calculation: High for counts of non-placeholder annotations in the supplied structured field; schema failures withhold the candidate. Interpretation: Low for real causes. This measures annotation coverage within a model-prompted six-category structure.

- The prompt asks for six categories. Their presence is partly a feature of the classification procedure.

- N/A placeholders are excluded. Annotation counts do not establish causal importance or verify the explanations.

- The diagram illustrates record 1; the aggregate coverage appears in a separate table.

IT-P002 - Information-security scores within the high-score group

Among the 32 records assigned an overall national-security score of at least 4 in the 2026-09-27 export, 17 have a score of at least 4 for information warfare and intelligence security, compared with 13 for societal stability and human rights. These are overlapping model-assigned categories within a small selected group, not estimates of national-security risk.

Chart. National-security radar. Radar example and supporting table.

Settings. All dates in the supplied export; partial-year records included only in this snapshot description; Choose an incident in the national-security listing and open its five-axis radar; ordinal scores 1-5, no counts/proportions toggle; stack or group by None; five separate impact ratings. Filters: incident_id: 544; overall_national_security_score_min: 4

Compatibility. The radar illustrates one record, not the aggregate count comparison. Use the attached five-category table for that comparison. The live mirror can differ from this export.

Confidence. Calculation: High for recorded scores, distinct IDs and threshold membership; missing scores never count as reaching a threshold. Interpretation: Limited to the supplied model labels. Small, selectively reported cohort. At threshold 5, 3 records qualify; the leading category or categories are Societal stability and human rights. Category ordering is threshold-dependent and must be reviewed for this export.

- Only 32 records meet the overall score threshold. Categories overlap and cannot be added.

- At threshold 5, three records qualify and societal stability leads. This ordering is threshold-dependent.

- The radar illustrates record 544; a separate table shows the 32-record comparison. Scores are model assessments, not measured national-security risk.

Published-insight audit

The captured statements below are assessed against the supplied export and their stated interpretation. Proposed numerical copy needs the September 27, 2026 export attribution. It should not be labeled a live chart reading unless the chart serves the same classified snapshot.

Unsupported

TIME-04 - The EU tier-1 annual increase

Captured wording. The proportion of incidents classified as level 1 (Unacceptable) under the EU AI Act increased each year since 2022, reaching 4% in 2025]

Assessment. 2022: 3/108 (2.8%); 2023: 9/179 (5.0%); 2024: 15/304 (4.9%); 2025: 24/452 (5.3%). Adjacent changes use unrounded fractions: 2023 to 2024 (-0.09 pp; within one record). Years before the claim's window (context only; small cells): 2016: 2/41 (4.9%); 2017: 2/50 (4.0%); 2018: 4/46 (8.7%); 2019: 6/44 (13.6%); 2020: 6/92 (6.5%); 2021: 7/80 (8.8%). 2022 is the lowest year of 2016-2025. 2026 year-to-date through 2026-09-27 (incomplete; context only, not comparable with whole years): 4/231 (1.7%). These are automated tier classifications, not legal rulings.

Proposed wording. Among incident records dated in each year, the share tagged EU AI Act tier 1, "Unacceptable," was 3 of 108 (2.8%) in 2022, 9 of 179 (5.0%) in 2023, 15 of 304 (4.9%) in 2024, and 24 of 452 (5.3%) in 2025. It did not increase every year. These automated tags are not legal determinations.

Display context. The numerical findings below use the September 27, 2026 Arcola export of 1,698 incident records. They describe this database and its automated classifications. Reporting, coverage, and classification changes can affect the patterns; these results do not estimate the prevalence of AI harm.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-timeline

Update needed

HOME-01 - The tracker's record count

Captured wording. The MIT AI Incident Tracker classifies 1,600 real-world report incidents from the AI Incident Database by risk, cause, harm, severity, and other relevant dimensions.

Assessment. The export contains 1,698 unique incident records. The static HTML fallback says 1,600; the latest rendered page shows 1,698. The live mirror nevertheless has different classifications. A static numerical claim should state its snapshot date or use the verified dynamic count.

Proposed wording. The MIT AI Incident Tracker classifies incident records from the AI Incident Database by risk, cause, harm, severity, and other dimensions.

Source: https://airisk.mit.edu/ai-incident-tracker

OVERVIEW-03 - The EU tier shares: 5.7% and 33%

Captured wording. 5.7% of reported incidents would be classed as level 1 (Unacceptable) and therefore prohibited under the EU AI Act. 33% would be classified as level 2 (High Risk)

Assessment. The current export gives 87/1,698 (5.1%) in tier 1 and 435/1,698 (25.6%) in tier 2. The published percentages are consistent with an earlier dataset composition, but confirming their origin requires the original snapshot. The reconstruction uses current labels and a conditional assumption about historical ID order, so it cannot establish the classifications in force in April 2025. Classifier tags do not establish legal prohibition.

Proposed wording. Across all dates in the export, 87 of 1,698 records (5.1%) are tagged EU AI Act tier 1, "Unacceptable," and 435 (25.6%) are tagged tier 2, "High Risk." These automated classifications do not determine whether a particular practice is legally prohibited.

Display context. The numerical findings below use the September 27, 2026 Arcola export of 1,698 incident records. They describe this database and its automated classifications. Reporting, coverage, and classification changes can affect the patterns; these results do not estimate the prevalence of AI harm.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-overview

Unresolved

HOME-02 - Incidents "on the rise"

Captured wording. AI incidents are on the rise, yet current databases struggle with inconsistent structure, limiting their utility for policymaking. The AI Incident Tracker project addresses this by creating a tool to classify AI incidents based on risks and harm severity. Raw reports from the AI Incident Database (AIID) are processed by Arcola AI's LLM pipeline, which categorizes them using established frameworks, such as the MIT Risk Repository and a harm severity rating system based on CSET's AI Harm Taxonomy.

Assessment. No time window or population exposure denominator is stated. Contextual complete-year record counts: 2021: 80/80 (100.0%); 2022: 108/108 (100.0%); 2023: 179/179 (100.0%); 2024: 304/304 (100.0%); 2025: 452/452 (100.0%). Primary system-safety domain: 378/1,698 (22.3%). Classifier missing-information flag: 42/1,698 (2.5%), which is not a measure of all reporting omissions.

Proposed wording. The AI Incident Tracker adds structure to reports in the AI Incident Database. Arcola AI's classification pipeline assigns risk and harm categories so researchers can examine the composition of reported incidents. Changes in database counts can reflect reporting and coverage as well as underlying events.

Source: https://airisk.mit.edu/ai-incident-tracker

HOME-06 - What the visualizations show

Captured wording. Key visualizations include bar charts and pie charts that display incident counts, proportions across domains (e.g., "System Failures," "Discrimination & Toxicity"), and trends in causal attributes. Additionally, insights highlight patterns such as the prevalence of system safety issues, intentional misuse trends, and incomplete reporting gaps.

Assessment. No time window or population exposure denominator is stated. Contextual complete-year record counts: 2021: 80/80 (100.0%); 2022: 108/108 (100.0%); 2023: 179/179 (100.0%); 2024: 304/304 (100.0%); 2025: 452/452 (100.0%). Primary system-safety domain: 378/1,698 (22.3%). Classifier missing-information flag: 42/1,698 (2.5%), which is not a measure of all reporting omissions.

Proposed wording. The charts show incident counts and distributions across risk domains, causal classifications, and harm assessments. Some views count multiple assignments per incident, so their shares describe assignments rather than distinct incidents. These patterns describe the database; they do not establish how common AI harms are across society.

Source: https://airisk.mit.edu/ai-incident-tracker

TIME-03 - Video, voice, and image generation

Captured wording. Numbers of reported incidents attributed to systems whose purpose is video, voice or image generation have increased dramatically since 2022, now totalling around half reported incidents.

Assessment. strict_primary: 2022: 19/108 (17.6%); 2023: 74/179 (41.3%); 2024: 147/304 (48.4%); 2025: 239/452 (52.9%); expanded_primary: 2022: 22/108 (20.4%); 2023: 75/179 (41.9%); 2024: 148/304 (48.7%); 2025: 243/452 (53.8%); strict_primary_secondary: 2022: 23/108 (21.3%); 2023: 77/179 (43.0%); 2024: 161/304 (53.0%); 2025: 256/452 (56.6%). All-date narrow share: 601/1,698 (35.4%); incomplete 2026: 101/231 (43.7%) (not a full-year comparison).

Proposed wording. Among incident records dated in each year, the number with a primary-purpose tag of Deepfake Video Generation, Deepfake Image Generation, Voice Generation, or Image Generation rose from 19 of 108 records (17.6%) in 2022 to 239 of 452 (52.9%) in 2025. These four tags define the group used in this comparison.

Display context. The numerical findings below use the September 27, 2026 Arcola export of 1,698 incident records. They describe this database and its automated classifications. Reporting, coverage, and classification changes can affect the patterns; these results do not estimate the prevalence of AI harm.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-timeline

TIME-06 - Post-deployment incidents

Captured wording. The overwhelming majority of incidents assessed were classified as post-deployment with no more than a few percent being pre-deployment in any year.

Assessment. Post-deployment overall: 1,663/1,698 (97.9%). Largest pre-deployment share across observed complete years: 1996 1/1 (100.0%). In the default chart's complete-year scope (2016-2025), maximum: 2016 4/41 (9.8%); full annual series in the component below. Sparse years and unstated scope/'few' definitions leave the parent wording ambiguous; see the literal-universe component. 2026 year-to-date (incomplete; context only): 8/231 (3.5%).

Proposed wording. Across all dates in the export, 1,663 of 1,698 records (97.9%) are tagged Post-deployment, 32 (1.9%) are tagged Pre-deployment, and 3 are tagged Other. This overall distribution does not imply the same proportions in every year.

Display context. The numerical findings below use the September 27, 2026 Arcola export of 1,698 incident records. They describe this database and its automated classifications. Reporting, coverage, and classification changes can affect the patterns; these results do not estimate the prevalence of AI harm. This all-date export summary covers a broader time span than the timeline's 2016-onward view. Present it with an all-date export table, or recalculate it for the displayed chart window and dataset.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-timeline

OVERVIEW-01 - "About half" tagged intentional

Captured wording. About half of all reported incidents were tagged as intentionally caused.

Assessment. All records: 830/1,698 (48.9%); omitting Other/unknown as a sensitivity: 830/1,696 (48.9%). No numerical tolerance is invented for 'About half'.

The measured share, 830 of 1,698 records (48.9%), is consistent with the ordinary meaning of "about half." The unresolved audit result reflects the lack of a specified tolerance, not a numerical contradiction. The proposed revision states the exact share and date.

Proposed wording. Across all dates in the export, 830 of 1,698 incident records (48.9%) are tagged Intentional.

Display context. The numerical findings below use the September 27, 2026 Arcola export of 1,698 incident records. They describe this database and its automated classifications. Reporting, coverage, and classification changes can affect the patterns; these results do not estimate the prevalence of AI harm.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-overview

OVERVIEW-06 - Fraud within malicious actors

Captured wording. Within domain 4 Malicious actors, the vast majority of reported incidents were in the subdomain '4.3 Fraud, scams and targeted manipulation'

Assessment. Primary within primary domain: 440/641 (68.6%); any subdomain within the same primary-domain population: 478/641 (74.6%); default bar floor: 438/639 (68.5%). 'Vast majority' has no declared threshold.

Proposed wording. Across all dates in the export, 440 of the 641 records whose primary domain is "4 Malicious actors" (68.6%) have the primary subdomain "4.3 Fraud, scams, and targeted manipulation."

Display context. The numerical findings below use the September 27, 2026 Arcola export of 1,698 incident records. They describe this database and its automated classifications. Reporting, coverage, and classification changes can affect the patterns; these results do not estimate the prevalence of AI harm. These all-date export totals cover a broader time span than the domain bar chart's 2016-onward view. Present them with an all-date export table, or recalculate them for the displayed chart window and dataset.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-overview

TIME-VALIDATION - Expert agreement: timeline page

Captured wording. The LLM classification tool has been developed iteratively, and its agreement with human expert consensus is comparable to the agreement between 2 independent human experts (article forthcoming). Spot-checks have been used to provide feedback on misclassifications and to iterate the tool, improving its reliability, however there are likely to remain incidents where the LLM classification does not match expert consensus.

Assessment. The June 2026 pilot is now published, but it tested ten incidents with earlier models and refined prompts on that same sample. It does not validate all current classifications. The published June pilot did not evaluate national-security scores. A separate shared preliminary study reports earlier national-security review, but it does not establish current-pipeline accuracy. The newly located ten-incident reference permits a limited comparison with current labels, shown under Pilot reference check.

Proposed wording. A June 2026 pilot compared eight models with consensus labels from two reviewers on 10 incidents. Some model and prompt combinations met the human agreement baseline on that sample. Prompt refinements used the same incidents, so these results require testing on a separate sample. The study does not establish the accuracy of all current classifications; the September 27 export predominantly uses a model not evaluated in that pilot.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-timeline

OVERVIEW-VALIDATION - Expert agreement: overview page

Captured wording. The LLM classification tool has been developed iteratively, and its agreement with human expert consensus is comparable to the agreement between 2 independent human experts (article forthcoming). Spot-checks have been used to provide feedback on misclassifications and to iterate the tool, improving its reliability, however there are likely to remain incidents where the LLM classification does not match expert consensus.

Assessment. The June 2026 pilot is now published, but it tested ten incidents with earlier models and refined prompts on that same sample. It does not validate all current classifications. The published June pilot did not evaluate national-security scores. A separate shared preliminary study reports earlier national-security review, but it does not establish current-pipeline accuracy. The newly located ten-incident reference permits a limited comparison with current labels, shown under Pilot reference check.

Proposed wording. A June 2026 pilot compared eight models with consensus labels from two reviewers on 10 incidents. Some model and prompt combinations met the human agreement baseline on that sample. Prompt refinements used the same incidents, so these results require testing on a separate sample. The study does not establish the accuracy of all current classifications; the September 27 export predominantly uses a model not evaluated in that pilot.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-overview

VIEW-VALIDATION - Expert agreement: incident view

Captured wording. The LLM classification tool has been developed iteratively, and its agreement with human expert consensus is comparable to the agreement between 2 independent human experts (article forthcoming). Spot-checks have been used to provide feedback on misclassifications and to iterate the tool, improving its reliability, however there are likely to remain incidents where the LLM classification does not match expert consensus.

Assessment. The June 2026 pilot is now published, but it tested ten incidents with earlier models and refined prompts on that same sample. It does not validate all current classifications. The published June pilot did not evaluate national-security scores. A separate shared preliminary study reports earlier national-security review, but it does not establish current-pipeline accuracy. The newly located ten-incident reference permits a limited comparison with current labels, shown under Pilot reference check.

Proposed wording. A June 2026 pilot compared eight models with consensus labels from two reviewers on 10 incidents. Some model and prompt combinations met the human agreement baseline on that sample. Prompt refinements used the same incidents, so these results require testing on a separate sample. The study does not establish the accuracy of all current classifications; the September 27 export predominantly uses a model not evaluated in that pilot.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-view

NATSEC-VALIDATION - Expert agreement: national-security view

Captured wording. The LLM classification tool has been developed iteratively, and its agreement with human expert consensus is comparable to the agreement between 2 independent human experts (article forthcoming). Spot-checks have been used to provide feedback on misclassifications and to iterate the tool, improving its reliability, however there are likely to remain incidents where the LLM classification does not match expert consensus.

Assessment. The June 2026 pilot is now published, but it tested ten incidents with earlier models and refined prompts on that same sample. It does not validate all current classifications. The published June pilot did not evaluate national-security scores. A separate shared preliminary study reports earlier national-security review, but it does not establish current-pipeline accuracy. The newly located ten-incident reference permits a limited comparison with current labels, shown under Pilot reference check.

Proposed wording. A June 2026 pilot compared eight models with consensus labels from two reviewers on 10 incidents. Some model and prompt combinations met the human agreement baseline on that sample. Prompt refinements used the same incidents, so these results require testing on a separate sample. The study does not establish the accuracy of all current classifications; the September 27 export predominantly uses a model not evaluated in that pilot. The pilot did not evaluate national-security impact scores.

Source: https://airisk.mit.edu/ai-incident-tracker/natsec-incident-view

Supported

TIME-01 - Fraud and scams over time

Captured wording. There has been a clear rise year on year in the proportion of incidents each year between 2021 and 2025 associated with 4.3 Fraud, scams and targeted manipulation

Assessment. Primary: 2021: 3/80 (3.8%); 2022: 13/108 (12.0%); 2023: 38/179 (21.2%); 2024: 105/304 (34.5%); 2025: 192/452 (42.5%). Any-label incident denominator: 2021: 6/80 (7.5%); 2022: 17/108 (15.7%); 2023: 56/179 (31.3%); 2024: 131/304 (43.1%); 2025: 227/452 (50.2%). Endpoint change +38.73 percentage points; counts change +189. 2026 year-to-date through 2026-09-27 (incomplete; context only, not comparable with whole years): 70/231 (30.3%).

Proposed wording. Among records dated 2021-2025, the share with the primary subdomain "4.3 Fraud, scams, and targeted manipulation" rose each year, from 3 of 80 (3.75%) in 2021 to 192 of 452 (42.48%) in 2025.

Display context. The numerical findings below use the September 27, 2026 Arcola export of 1,698 incident records. They describe this database and its automated classifications. Reporting, coverage, and classification changes can affect the patterns; these results do not estimate the prevalence of AI harm.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-timeline

TIME-02 - Capability and robustness over time

Captured wording. The proportion of reported incidents attributed to lack of capability or robustness has been decreasing between 2021 and 2025

Assessment. Primary: 2021: 33/80 (41.2%); 2022: 33/108 (30.6%); 2023: 37/179 (20.7%); 2024: 44/304 (14.5%); 2025: 45/452 (10.0%). Any-label incident denominator: 2021: 53/80 (66.2%); 2022: 62/108 (57.4%); 2023: 72/179 (40.2%); 2024: 108/304 (35.5%); 2025: 121/452 (26.8%). Endpoint change -31.29 percentage points; counts change +12. 2026 year-to-date through 2026-09-27 (incomplete; context only, not comparable with whole years): 43/231 (18.6%).

Proposed wording. Among records dated 2021-2025, the share with the primary subdomain "7.3 Lack of capability or robustness" fell each year, from 33 of 80 (41.25%) to 45 of 452 (9.96%). The count increased, so the falling share does not establish improved AI robustness.

Display context. The numerical findings below use the September 27, 2026 Arcola export of 1,698 incident records. They describe this database and its automated classifications. Reporting, coverage, and classification changes can affect the patterns; these results do not estimate the prevalence of AI harm.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-timeline

TIME-05 - Intentional and human-caused incidents

Captured wording. The proportion of intentionally caused incidents has increased each year from 2021-2025, as has the proportion attributed to Human rather than AI entities.

Assessment. causal_intent: 2021: 17/80 (21.2%); 2022: 36/108 (33.3%); 2023: 90/179 (50.3%); 2024: 192/304 (63.2%); 2025: 291/452 (64.4%). causal_entity: 2021: 22/80 (27.5%); 2022: 41/108 (38.0%); 2023: 94/179 (52.5%); 2024: 198/304 (65.1%); 2025: 310/452 (68.6%). 2026 year-to-date through 2026-09-27 (incomplete; context only, not comparable with whole years): 130/231 (56.3%); 2026 year-to-date through 2026-09-27 (incomplete; context only, not comparable with whole years): 142/231 (61.5%).

Proposed wording. Among records dated 2021-2025, the share tagged Intentional rose each year, from 17 of 80 (21.25%) to 291 of 452 (64.38%). The share whose causal entity was tagged Human also rose each year, from 22 of 80 (27.5%) to 310 of 452 (68.58%). These tags are model assessments of causation.

Display context. The numerical findings below use the September 27, 2026 Arcola export of 1,698 incident records. They describe this database and its automated classifications. Reporting, coverage, and classification changes can affect the patterns; these results do not estimate the prevalence of AI harm.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-timeline

OVERVIEW-02 - Pre- and post-deployment shares

Captured wording. Almost all reported incidents have been classified as post-deployment, with only 2% pre-deployment

Assessment. Pre: 32/1,698 (1.9%); post: 1,663/1,698 (97.9%); Other: 3/1,698 (0.2%).

Proposed wording. Across all dates in the export, 1,663 of 1,698 records (97.9%) are tagged Post-deployment, 32 (1.9%) are tagged Pre-deployment, and 3 are tagged Other.

Display context. The numerical findings below use the September 27, 2026 Arcola export of 1,698 incident records. They describe this database and its automated classifications. Reporting, coverage, and classification changes can affect the patterns; these results do not estimate the prevalence of AI harm.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-overview

OVERVIEW-04 - The largest primary-purpose category

Captured wording. The primary purpose category attributed to the highest number of incidents is Deepfake Video Generation

Assessment. Deepfake Video Generation: 316/1,698 (18.6%); Chatbot: 181/1,698 (10.7%); Deepfake Image Generation: 135/1,698 (8.0%). Ranking counts primary labels over all incident dates, including the partial current year as snapshot composition.

Proposed wording. Deepfake Video Generation is the most frequent single primary-purpose tag, assigned to 316 of 1,698 records (18.6%), followed by Chatbot with 181 (10.7%). The chart's "Other" slice combines multiple less frequent purpose tags.

Display context. The numerical findings below use the September 27, 2026 Arcola export of 1,698 incident records. They describe this database and its automated classifications. Reporting, coverage, and classification changes can affect the patterns; these results do not estimate the prevalence of AI harm.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-overview

OVERVIEW-05 - The largest primary risk domain

Captured wording. The domain with most reported incidents across the full dataset is '4 Malicious actors' followed by '7 AI system safety, failures, & limitations'

Assessment. 4 Malicious actors: 641/1,698 (37.8%); 7 AI system safety, failures, & limitations: 378/1,698 (22.3%); 1 Discrimination & Toxicity: 281/1,698 (16.5%). Ranking counts primary labels over all incident dates, including the partial current year as snapshot composition.

Proposed wording. Across all dates in the export, the most frequent primary domain is "4 Malicious actors," with 641 of 1,698 records (37.8%), followed by "7 AI system safety, failures, & limitations," with 378 (22.3%).

Display context. The numerical findings below use the September 27, 2026 Arcola export of 1,698 incident records. They describe this database and its automated classifications. Reporting, coverage, and classification changes can affect the patterns; these results do not estimate the prevalence of AI harm. These all-date export totals cover a broader time span than the domain bar chart's 2016-onward view. Present them with an all-date export table, or recalculate them for the displayed chart window and dataset.

Source: https://airisk.mit.edu/ai-incident-tracker/incident-overview

Additional findings and prepared corrections

Severity headline. The public headline labeled Highest Direct Severity Score includes direct, indirect, and inferred values across thirteen dimensions. Record 1702 supplies a counterexample: direct and indirect scores do not exceed 2, while the headline is 3. A patch changes the label to Highest assessed severity and explains the included channels. It applies to the captured source; compilation in the maintainer's application remains to be done.

Severity preset. The old minSeverity=4 parameter returned the same 1,627 records as no threshold for the checked timeline query. The corrected minSeverityHighest=4 returned 348. Safari interactions corroborated the ineffective preset. Corrected timeline and national-security presets are prepared; the national-security level-5 label should be Critical.

Record 1607. The source article for Bill Oliver's New Brunswick legislative speech includes unrelated U.S. lawyer examples. Several classifications incorrectly absorb those examples. In the earlier 2,305-measure run, excluding the record changed 494 measures, mainly denominators and current-year sensitivities; complete-year headline findings are unchanged. This targeted source check cannot estimate a dataset-wide error rate. Arcola can reclassify the target incident and recompute its dependent fields.

Definitions. All 69 zero-coded severity cells have N/A in their matching narrative and occur in two records, but the exporter's convention remains unconfirmed. The public scale covers ten dimensions, not the added psychological, epistemic, and CSEA thresholds. An older behavior rubric was recovered; its applicability to the current classifier remains unconfirmed. National-security autonomy definitions are documented: human-controlled systems assist human decisions; human-supervised systems act autonomously under oversight; full-autonomy systems make and execute decisions without intervention.

Validation. The published June 2026 pilot used 10 incidents, two human reviewers, and eight models, with refinement on the same sample. Neither model label in the current export appears in that pilot's tested list. A held-out current-pipeline evaluation is still needed. National-security impact scores were not evaluated.

Snapshot differences. The live mirror and export both contained 1,698 records during checks, but the live purpose pie showed 311 deepfake-video records (18.3%) versus 316 (18.6%) in the export; live EU tier 1 was 5.6% versus 5.1%. The source and date must accompany numerical claims.

Methods and recurring operation

Discovery. The declared process searched 424 comparisons, selected six additional candidates, and retained five audited baseline examples. It considered support, effect, stability, and chart fit, with an emergence rule for a small starting count. Individual weight changes of 0.1 did not change the selected set. The shortlist is exploratory, not a statistical-significance ranking.

Verification. The expanded run contains 2,340 calculations and 4,680 numerator/denominator ID references. Independent code reproduced all 35 new profile calculations, 70 new ID-set comparisons, and changes across 49 historical fields. Two earlier independent implementations reproduced the 1,486 inherited calculations. v1 and historical runs remain preserved. Real drafting and separate model review covered the earlier eleven-candidate set. The two new profiles have session review and independent code checks, but no new paid model review.

Chart coverage. All seven chart views now have candidate coverage. Individual radars and the causes diagram illustrate specific records; supporting tables carry aggregate comparisons. The new descriptive profiles do not relax the 50-record rule for comparative trend findings or establish causal effects.

Browser coverage. Representative public-site interactions were exercised in the in-app browser and Safari. This was not an exhaustive filter, incident, browser, or operating-system matrix. The redesigned local dashboard has content and link checks; its browser rendering was unavailable in this environment.

Monthly rehearsal. The unchanged-input run checked 59 registered entries across 10 sources: 52 quotations remained present and seven application-source entries were unchanged. It made no model calls and performed no recalculation. Failed or unfinished reviews remain visible even when the input is unchanged.

Schedule. A local check is configured for the 27th at 9 a.m. New York time. It requires the Mac, Codex, authenticated model access, and an available CSV. Changed-input runs can use metered models; operating accounts and billing ownership must remain explicit. The scheduled check now retrieves eligible dated files from the shared Google Drive tracker folders. Arcola's upstream generation and delivery of new exports has not been established.

Remaining work and ownership

Arcola data delivery. A genuine 1,200-record historical comparison is complete: 498 added IDs and no removed IDs, with changes tracked across 49 mapped fields. A second export verified from the current Arcola pipeline and upstream regular delivery remain required; scheduled Drive retrieval is configured.

Arcola and taxonomy owners. Confirm zero semantics, the current version of the located older behavior rubric, and the added severity rubrics; correct record 1607. We can update affected calculations and wording.

Arcola and qualified reviewers. The ten-incident pilot-reference diagnostic is complete. Current labels match 67/100 harm ratings, 7/10 primary domains, 5/10 EU tiers, 4/10 causal-entity labels, 9/10 intent labels and 10/10 timing labels. These are diagnostic results, not population accuracy. A representative held-out evaluation still requires independent expert judgments.

Historical source owner. Supply the original classified snapshot if the origin of 5.7% and 33% needs confirmation. Current-label subsets do not prove the historical explanation.

Site maintainer. Provide the application build and CMS access, approve the copy, and apply and test the prepared patch and presets before publication.